

1: Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from scipy.stats import skew, kurtosis

from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

from statsmodels.tools.sm_exceptions import ValueWarning

from tensorflow.keras import backend as K
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout, GroupNormalization
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import load_model
from tensorflow.keras.regularizers import l2
from tensorflow.keras.losses import Huber
from tcn import TCN

import shap
import gc
import joblib
import warnings
warnings.filterwarnings("ignore", category=UserWarning)
warnings.filterwarnings("ignore", category=FutureWarning)
warnings.filterwarnings("ignore", category=ValueWarning)

import os
import sys
print(sys.executable)
```

b:\DCU\Practicum\Proj\App\venv_3_11\Scripts\python.exe

b:\DCU\Practicum\Proj\App\venv_3_11\Lib\site-packages\tqdm\auto.py:21: TqdmWarning: IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm

2: Load and Explore Dataset

```
In [2]: multimodal = pd.read_csv("train_dataset.csv", parse_dates=["Date"])
multimodal.drop(columns=['Next_Close', 'Next_3_Close', 'Next_7_Close'], inp
lace=True)
multimodal.head()
```

Out[2]:

	Date	Open	High	Low	Close	Volume	Adj Close	I
0	2008-08-08	11432.089844	11759.959961	11388.040039	11734.320312	212830000	11734.320312	
1	2008-08-11	11729.669922	11867.110352	11675.530273	11782.349609	183190000	11782.349609	
2	2008-08-12	11781.700195	11782.349609	11601.519531	11642.469727	173590000	11642.469727	
3	2008-08-13	11632.809570	11633.780273	11453.339844	11532.959961	182550000	11532.959961	
4	2008-08-14	11532.070312	11718.280273	11450.889648	11615.929688	159790000	11615.929688	

5 rows × 53 columns



```
In [3]: multimodal.shape
```


Out[3]: (1591, 53)

```
In [4]: missing_values = multimodal.isnull().sum()
print("\nMissing values per column:")
missing_values[missing_values > 0]
```

Missing values per column:

Out[4]: Series([], dtype: int64)

```
In [5]: plt.figure(figsize=(14, 7))
plt.plot(pd.to_datetime(multimodal['Date']), multimodal['Close'])
plt.title('Stock Closing Price Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.xticks(rotation=45)
plt.tight_layout()
plt.savefig("B:/DCU/Practicum/Proj/Outputs/closing_price_plot.png")
plt.show()
```



3: Data Preprocessing

```
In [6]: df_targets = multimodal[["Date", "Target", "Label", "Close"]].copy()

df_arima = multimodal[["Date", "Close"]].copy()
df_arima.set_index("Date", inplace=True)

drop_cols_lstm = ["Date", "Label", "Target"]
df_lstm = multimodal.drop(columns=drop_cols_lstm).copy()

drop_cols_tcn = ["Date", "Label", "Target"]
df_tcn = multimodal.drop(columns=drop_cols_tcn).copy()
```

```
In [7]: # Function to clear TensorFlow memory before running any model so that the
model doesnt predict on cached data
def clear_tf_memory():
    K.clear_session()
    gc.collect()
```

4: ARIMA Model


```
In [8]: result = adfuller(df_arima["Close"])
print(f"ADF Statistic: {result[0]}")
print(f"p-value: {result[1]}")
```

ADF Statistic: 0.49003049114940356
p-value: 0.9845669104126585

```
In [9]: plt.figure(figsize=(14, 5))
plt.subplot(1, 2, 1)
plot_acf(df_arima["Close"], lags=40, ax=plt.gca())
plt.title("ACF: Close")

plt.subplot(1, 2, 2)
plot_pacf(df_arima["Close"], lags=40, ax=plt.gca(), method='ywmm')
plt.title("PACF: Close")

plt.tight_layout()
plt.savefig("B:/DCU/Practicum/Proj/Outputs/arima_acf_pacf.png")
plt.show()
```



```

In [10]: def arima_true_forecast(df_arima, forecast_horizon=1, order=(1, 1, 1), save
_dir=None, plot=True):
    print(df_arima.columns)
    close_series = df_arima['Close'].values
    history = close_series[:-(forecast_horizon + 1)].tolist()
    y_true = []
    y_pred = []

    for t in range(len(close_series) - forecast_horizon):
        try:
            model = ARIMA(history, order=order).fit()
            forecast = model.forecast(steps=forecast_horizon)
            y_pred.append(forecast[-1])
            y_true.append(close_series[t + forecast_horizon])
        except:
            y_pred.append(np.nan)
            y_true.append(np.nan)
        history.append(close_series[t])

    y_true = np.array(y_true)
    y_pred = np.array(y_pred)
    mask = ~np.isnan(y_pred)
    y_true = y_true[mask]
    y_pred = y_pred[mask]

    date_series = df_arima.index[forecast_horizon : forecast_horizon + len
(y_true)]

    # Metrics
    r2 = r2_score(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    mae = mean_absolute_error(y_true, y_pred)

    print(f"ARIMA Forecast Horizon = {forecast_horizon}")
    print(f"R2 = {r2:.4f}, RMSE = {rmse:.2f}, MAE = {mae:.2f}")

    if plot:
        plt.figure(figsize=(14, 5))
        plt.plot(date_series, y_true, label='Actual', color='blue')
        plt.plot(date_series, y_pred, label='Forecast', color='orange')
        plt.title(f"ARIMA True Forecast (t+{forecast_horizon})")
        plt.xlabel("Date")
        plt.ylabel("Close Price")
        plt.legend()
        plt.tight_layout()
        plt.savefig(f"B:/DCU/Practicum/Proj/Outputs/arima_train_t_plus_{for
ecast_horizon}.png")
        plt.show()

    if save_dir:
        os.makedirs(save_dir, exist_ok=True)
        model = ARIMA(close_series.tolist(), order=order).fit()
        save_path = os.path.join(save_dir, f'arima_t_plus_{forecast_horizo
n}.pkl')
        joblib.dump(model, save_path)
        print(f"Model saved to {save_path}")

    return {
        "horizon": forecast_horizon,
        "r2": r2,

```


```
        "rmse": rmse,  
        "mae": mae  
    }  
  
results = []  
for h in [1, 3, 7]:  
    clear_tf_memory()  
    res = arima_true_forecast(df_arima, forecast_horizon=h, order=(1,1,1),  
save_dir="B:/DCU/Practicum/Proj/Models")  
    results.append(res)
```

WARNING:tensorflow:From b:\DCU\Practicum\Proj\AppData\Local\Temp\venv_3_11\Lib\site-packages\keras\src\backend\tensorflow_backend.py:82: The name tf.reset_default_graph is deprecated. Please use tf.compat.v1.reset_default_graph instead.

Index(['Close'], dtype='object')

ARIMA Forecast Horizon = 1

$R^2 = 0.9921$, RMSE = 238.87, MAE = 136.61




Model saved to B:/DCU/Practicum/Proj/Models\arima_t_plus_1.pkl

Index(['Close'], dtype='object')

ARIMA Forecast Horizon = 3

$R^2 = 0.9877$, RMSE = 297.54, MAE = 190.21




Model saved to B:/DCU/Practicum/Proj/Models\arima_t_plus_3.pkl

Index(['Close'], dtype='object')

ARIMA Forecast Horizon = 7

$R^2 = 0.9800$, RMSE = 380.06, MAE = 260.24



Model saved to B:/DCU/Practicum/Proj/Models\arima_t_plus_7.pkl

LSTM Model

```
In [13]: def train_lstm_multistep(df_lstm, forecast_horizon=1, window_size=60, epoch
s=50, batch_size=32, stacked=False):
    print(f"\nTraining LSTM ({'Stacked' if stacked else 'Single'}) model fo
r horizon t+{forecast_horizon}")

    close_values = df_lstm['Close'].values

    # Scale input and output
    X_scaler = MinMaxScaler()
    y_scaler = MinMaxScaler()
    X_scaled = X_scaler.fit_transform(df_lstm.values)
    y_scaled = y_scaler.fit_transform(close_values.reshape(-1, 1))

    # Create sequences
    X_seq, y_seq = [], []
    for i in range(window_size, len(X_scaled) - forecast_horizon):
        X_seq.append(X_scaled[i - window_size:i])
        y_seq.append(y_scaled[i + forecast_horizon])
    X_seq, y_seq = np.array(X_seq), np.array(y_seq)

    X_train, y_train = X_seq, y_seq

    # Build model
    model = Sequential()

    if stacked:
        reg = l2(1e-5)
        model.add(LSTM(64, return_sequences=True, input_shape=(X_train.shap
e[1], X_train.shape[2]),
                      kernel_regularizer=reg, recurrent_regularizer=reg))
        model.add(Dropout(0.3))
        model.add(LSTM(32, kernel_regularizer=reg, recurrent_regularizer=re
g))
        model.add(Dropout(0.3))
    else:
        model.add(LSTM(64, input_shape=(X_train.shape[1], X_train.shape
[2])))
        model.add(Dropout(0.1))
        model.add(Dense(1))

    optimizer = Adam(learning_rate=1e-4)
    model.compile(optimizer=optimizer, loss='mse')
    early_stop = EarlyStopping(monitor='val_loss', patience=5, restore_best
_weights=True)

    model.fit(
        X_train, y_train,
        epochs=epochs,
        batch_size=batch_size,
        validation_split=0.1,
        callbacks=[early_stop],
        verbose=1
    )

    # Save model and scalers
    model_name = f"lstm_tplus{forecast_horizon}{'_stacked' if stacked else
```



```

'_simple'}"
    base_path = "B:/DCU/Practicum/Proj/Models"
    os.makedirs(base_path, exist_ok=True)

    model.save(f"{base_path}/{model_name}.keras")
    joblib.dump(X_scaler, f"{base_path}/{model_name}_scalerX.pkl")
    joblib.dump(y_scaler, f"{base_path}/{model_name}_scalerY.pkl")

    # Save feature list
    feature_columns = df_lstm.columns.tolist()
    joblib.dump(feature_columns, f"{base_path}/{model_name}_features.pkl")
    print(f"Saved model and scalers: {model_name}")

    # Predict on train set to evaluate
    y_pred_scaled = model.predict(X_train)
    y_pred = y_scaler.inverse_transform(y_pred_scaled)
    y_true = y_scaler.inverse_transform(y_train.reshape(-1, 1))

    date_series = multimodal['Date'].iloc[window_size + forecast_horizon :
window_size + forecast_horizon + len(y_pred)]

    # Metrics
    r2 = r2_score(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    mae = mean_absolute_error(y_true, y_pred)
    print(f"Train R²: {r2:.4f}, RMSE: {rmse:.2f}, MAE: {mae:.2f}")

    # Plot
    plt.figure(figsize=(14, 5))
    plt.plot(date_series, y_true, label='Actual', color='blue')
    plt.plot(date_series, y_pred, label='Predicted', color='orange')
    plt.title(f"LSTM ({'Stacked' if stacked else 'Single'}) Forecast (t+{fo
recast_horizon}) - Training Set")
    plt.xlabel("Date")
    plt.ylabel("Close Price")
    plt.legend()
    plt.tight_layout()
    plt.savefig(f"B:/DCU/Practicum/Proj/Outputs/lstm_train_t_plus_{forecast
_horizon}{'_stacked' if stacked else '_simple'}.png")
    plt.show()

    return model, {"r2": r2, "rmse": rmse, "mae": mae}


























for horizon in [1, 3, 7]:
    clear_tf_memory()
    train_lstm_multistep(df_lstm, forecast_horizon=horizon, stacked=False)

for horizon in [1, 3, 7]:
    clear_tf_memory()
    train_lstm_multistep(df_lstm, forecast_horizon=horizon, stacked=True)


```

Training LSTM (Single) model for horizon t+1

```

Epoch 1/50
44/44  2s 20ms/step - loss: 0.0427 - val_loss: 0.0370
Epoch 2/50
44/44  1s 17ms/step - loss: 0.0146 - val_loss: 0.0056
Epoch 3/50
44/44  1s 15ms/step - loss: 0.0089 - val_loss: 0.0039
Epoch 4/50
44/44  1s 16ms/step - loss: 0.0084 - val_loss: 0.0033
Epoch 5/50
44/44  1s 17ms/step - loss: 0.0061 - val_loss: 0.0031
Epoch 6/50
44/44  1s 17ms/step - loss: 0.0058 - val_loss: 0.0016
Epoch 7/50
44/44  1s 17ms/step - loss: 0.0053 - val_loss: 0.0014
Epoch 8/50
44/44  1s 14ms/step - loss: 0.0048 - val_loss: 0.0016
Epoch 9/50
44/44  1s 15ms/step - loss: 0.0039 - val_loss: 0.0013
Epoch 10/50
44/44  1s 15ms/step - loss: 0.0039 - val_loss: 0.0017
Epoch 11/50
44/44  1s 17ms/step - loss: 0.0037 - val_loss: 0.0017
Epoch 12/50
44/44  1s 18ms/step - loss: 0.0037 - val_loss: 0.0016
Epoch 13/50
44/44  1s 17ms/step - loss: 0.0031 - val_loss: 0.0010
Epoch 14/50
44/44  1s 19ms/step - loss: 0.0028 - val_loss: 0.0016
Epoch 15/50
44/44  1s 15ms/step - loss: 0.0029 - val_loss: 0.0022
Epoch 16/50
44/44  1s 17ms/step - loss: 0.0029 - val_loss: 0.0020
Epoch 17/50
44/44  1s 16ms/step - loss: 0.0026 - val_loss: 9.3109e-04
Epoch 18/50
44/44  1s 16ms/step - loss: 0.0026 - val_loss: 0.0011
Epoch 19/50
44/44  1s 16ms/step - loss: 0.0022 - val_loss: 6.5980e-04
Epoch 20/50
44/44  1s 17ms/step - loss: 0.0021 - val_loss: 0.0015
Epoch 21/50
44/44  1s 16ms/step - loss: 0.0020 - val_loss: 8.1791e-04
Epoch 22/50
44/44  1s 15ms/step - loss: 0.0019 - val_loss: 8.4484e-04
Epoch 23/50
44/44  1s 14ms/step - loss: 0.0020 - val_loss: 0.0012
Epoch 24/50
44/44  1s 14ms/step - loss: 0.0019 - val_loss: 8.5470e-04
Saved model and scalers: lstm_tplus1_simple
48/48  0s 7ms/step
Train R2: 0.9881, RMSE: 294.14, MAE: 225.36

```



Training LSTM (Single) model for horizon t+3

Epoch 1/50

43/43 ————— 2s 21ms/step - loss: 0.0490 - val_loss: 0.0317

Epoch 2/50

43/43 ————— 1s 15ms/step - loss: 0.0130 - val_loss: 0.0048

Epoch 3/50

43/43 ————— 1s 15ms/step - loss: 0.0119 - val_loss: 0.0041

Epoch 4/50

43/43 ————— 1s 15ms/step - loss: 0.0100 - val_loss: 0.0033

Epoch 5/50

43/43 ————— 1s 15ms/step - loss: 0.0080 - val_loss: 0.0035

Epoch 6/50

43/43 ————— 1s 15ms/step - loss: 0.0071 - val_loss: 0.0040

Epoch 7/50

43/43 ————— 1s 15ms/step - loss: 0.0066 - val_loss: 0.0021

Epoch 8/50

43/43 ————— 1s 15ms/step - loss: 0.0056 - val_loss: 0.0011

Epoch 9/50

43/43 ————— 1s 15ms/step - loss: 0.0052 - val_loss: 9.8040e-04

Epoch 10/50

43/43 ————— 1s 15ms/step - loss: 0.0046 - val_loss: 9.2705e-04

Epoch 11/50

43/43 ————— 1s 15ms/step - loss: 0.0042 - val_loss: 9.5220e-04

Epoch 12/50

43/43 ————— 1s 15ms/step - loss: 0.0039 - val_loss: 0.0013

Epoch 13/50

43/43 ————— 1s 16ms/step - loss: 0.0040 - val_loss: 9.4911e-04

Epoch 14/50

43/43 ————— 1s 16ms/step - loss: 0.0037 - val_loss: 0.0011


Epoch 15/50

43/43 ————— 1s 15ms/step - loss: 0.0033 - val_loss: 0.0012

Saved model and scalers: lstm_tplus3_simple





























48/48 ————— 0s 7ms/step

Train R²: 0.9863, RMSE: 315.44, MAE: 241.68



Training LSTM (Single) model for horizon t+7

```

Epoch 1/50
43/43  2s 17ms/step - loss: 0.0637 - val_loss: 0.0921
Epoch 2/50
43/43  1s 14ms/step - loss: 0.0210 - val_loss: 0.0226
Epoch 3/50
43/43  1s 13ms/step - loss: 0.0087 - val_loss: 0.0042
Epoch 4/50
43/43  1s 13ms/step - loss: 0.0072 - val_loss: 0.0046
Epoch 5/50
43/43  1s 13ms/step - loss: 0.0060 - val_loss: 0.0033
Epoch 6/50
43/43  1s 13ms/step - loss: 0.0058 - val_loss: 0.0034
Epoch 7/50
43/43  1s 13ms/step - loss: 0.0048 - val_loss: 0.0053
Epoch 8/50
43/43  1s 13ms/step - loss: 0.0045 - val_loss: 0.0034
Epoch 9/50
43/43  1s 13ms/step - loss: 0.0041 - val_loss: 0.0031
Epoch 10/50
43/43  1s 13ms/step - loss: 0.0044 - val_loss: 0.0028
Epoch 11/50
43/43  1s 13ms/step - loss: 0.0037 - val_loss: 0.0031
Epoch 12/50
43/43  1s 13ms/step - loss: 0.0037 - val_loss: 0.0022
Epoch 13/50
43/43  1s 13ms/step - loss: 0.0034 - val_loss: 0.0016
Epoch 14/50
43/43  1s 13ms/step - loss: 0.0031 - val_loss: 0.0012
Epoch 15/50
43/43  1s 13ms/step - loss: 0.0030 - val_loss: 0.0015
Epoch 16/50
43/43  1s 13ms/step - loss: 0.0029 - val_loss: 0.0012
Epoch 17/50
43/43  1s 13ms/step - loss: 0.0029 - val_loss: 0.0011
Epoch 18/50
43/43  1s 13ms/step - loss: 0.0027 - val_loss: 0.0013
Epoch 19/50
43/43  1s 13ms/step - loss: 0.0026 - val_loss: 0.0017
Epoch 20/50
43/43  1s 13ms/step - loss: 0.0027 - val_loss: 0.0014
Epoch 21/50
43/43  1s 13ms/step - loss: 0.0024 - val_loss: 0.0011
Epoch 22/50
43/43  1s 13ms/step - loss: 0.0023 - val_loss: 9.6655e-04
Epoch 23/50
43/43  1s 13ms/step - loss: 0.0024 - val_loss: 8.4553e-04
Epoch 24/50
43/43  1s 14ms/step - loss: 0.0022 - val_loss: 0.0013
Epoch 25/50
43/43  1s 13ms/step - loss: 0.0022 - val_loss: 8.9481e-04
Epoch 26/50
43/43  1s 13ms/step - loss: 0.0022 - val_loss: 8.6509e-04
Epoch 27/50
43/43  1s 13ms/step - loss: 0.0018 - val_loss: 0.0012
Epoch 28/50
43/43  1s 13ms/step - loss: 0.0020 - val_loss: 8.1267e-04

```

-04

Epoch 29/50

43/43  **1s** 13ms/step - loss: 0.0019 - val_loss: 7.1048e


-04

Epoch 30/50

43/43  **1s** 13ms/step - loss: 0.0020 - val_loss: 7.2605e

-04

Epoch 31/50

43/43  **1s** 13ms/step - loss: 0.0020 - val_loss: 0.0011

Epoch 32/50

43/43  **1s** 13ms/step - loss: 0.0021 - val_loss: 8.5533e

-04

Epoch 33/50

43/43  **1s** 13ms/step - loss: 0.0019 - val_loss: 7.5960e

-04


Epoch 34/50

43/43  **1s** 13ms/step - loss: 0.0019 - val_loss: 0.0016

Saved model and scalers: lstm_tplus7_simple

48/48  **0s** 7ms/step

Train R^2 : 0.9855, RMSE: 323.81, MAE: 246.69




Training LSTM (Stacked) model for horizon t+1

Epoch 1/50

44/44  3s 28ms/step - loss: 0.1739 - val_loss: 0.0531


Epoch 2/50

44/44  1s 23ms/step - loss: 0.0360 - val_loss: 0.0362

Epoch 3/50

44/44  1s 23ms/step - loss: 0.0209 - val_loss: 0.0116

Epoch 4/50

44/44  1s 22ms/step - loss: 0.0158 - val_loss: 0.0086

Epoch 5/50

44/44  1s 22ms/step - loss: 0.0150 - val_loss: 0.0080

Epoch 6/50

44/44  1s 22ms/step - loss: 0.0134 - val_loss: 0.0057

Epoch 7/50

44/44  1s 23ms/step - loss: 0.0130 - val_loss: 0.0074

Epoch 8/50

44/44  1s 22ms/step - loss: 0.0113 - val_loss: 0.0058

Epoch 9/50

44/44  1s 22ms/step - loss: 0.0118 - val_loss: 0.0052

Epoch 10/50

44/44  1s 23ms/step - loss: 0.0118 - val_loss: 0.0067

Epoch 11/50

44/44  1s 22ms/step - loss: 0.0106 - val_loss: 0.0085

Epoch 12/50

44/44  1s 23ms/step - loss: 0.0114 - val_loss: 0.0054

Epoch 13/50

44/44  1s 23ms/step - loss: 0.0104 - val_loss: 0.0049

Epoch 14/50

44/44  1s 23ms/step - loss: 0.0111 - val_loss: 0.0051

Epoch 15/50

44/44  1s 23ms/step - loss: 0.0102 - val_loss: 0.0046

Epoch 16/50

44/44  1s 23ms/step - loss: 0.0097 - val_loss: 0.0055

Epoch 17/50

44/44  1s 22ms/step - loss: 0.0096 - val_loss: 0.0051

Epoch 18/50

44/44  1s 23ms/step - loss: 0.0094 - val_loss: 0.0065

Epoch 19/50

44/44  1s 23ms/step - loss: 0.0095 - val_loss: 0.0052


Epoch 20/50

44/44  1s 23ms/step - loss: 0.0091 - val_loss: 0.0059

Saved model and scalers: lstm_tplus1_stacked

48/48  1s 11ms/step

Train R^2 : 0.9826, RMSE: 355.83, MAE: 272.09



Training LSTM (Stacked) model for horizon t+3

Epoch 1/50

43/43  3s 29ms/step - loss: 0.0749 - val_loss: 0.0647

Epoch 2/50

43/43  1s 23ms/step - loss: 0.0342 - val_loss: 0.0112

Epoch 3/50

43/43  1s 23ms/step - loss: 0.0212 - val_loss: 0.0064

Epoch 4/50

43/43  1s 23ms/step - loss: 0.0176 - val_loss: 0.0072

Epoch 5/50

43/43  1s 23ms/step - loss: 0.0149 - val_loss: 0.0052

Epoch 6/50

43/43  1s 23ms/step - loss: 0.0135 - val_loss: 0.0051

Epoch 7/50

43/43  1s 23ms/step - loss: 0.0125 - val_loss: 0.0038

Epoch 8/50

43/43  1s 23ms/step - loss: 0.0132 - val_loss: 0.0035

Epoch 9/50

43/43  1s 23ms/step - loss: 0.0118 - val_loss: 0.0047

Epoch 10/50

43/43  1s 23ms/step - loss: 0.0108 - val_loss: 0.0038

Epoch 11/50

43/43  1s 23ms/step - loss: 0.0112 - val_loss: 0.0048

Epoch 12/50

43/43  1s 23ms/step - loss: 0.0107 - val_loss: 0.0038


Epoch 13/50

43/43  1s 23ms/step - loss: 0.0097 - val_loss: 0.0037

Saved model and scalers: lstm_tplus3_stacked

48/48  1s 11ms/step

Train R^2 : 0.9858, RMSE: 320.75, MAE: 251.53



Training LSTM (Stacked) model for horizon t+7

Epoch 1/50

43/43 ————— 3s 29ms/step - loss: 0.0723 - val_loss: 0.0758

Epoch 2/50

43/43 ————— 1s 23ms/step - loss: 0.0271 - val_loss: 0.0129

Epoch 3/50

43/43 ————— 1s 23ms/step - loss: 0.0178 - val_loss: 0.0084

Epoch 4/50

43/43 ————— 1s 23ms/step - loss: 0.0166 - val_loss: 0.0050

Epoch 5/50

43/43 ————— 1s 23ms/step - loss: 0.0141 - val_loss: 0.0062

Epoch 6/50

43/43 ————— 1s 23ms/step - loss: 0.0149 - val_loss: 0.0042

Epoch 7/50

43/43 ————— 1s 23ms/step - loss: 0.0126 - val_loss: 0.0066

Epoch 8/50

43/43 ————— 1s 23ms/step - loss: 0.0132 - val_loss: 0.0043

Epoch 9/50

43/43 ————— 1s 23ms/step - loss: 0.0122 - val_loss: 0.0043

Epoch 10/50

43/43 ————— 1s 24ms/step - loss: 0.0118 - val_loss: 0.0059


Epoch 11/50

43/43 ————— 1s 23ms/step - loss: 0.0116 - val_loss: 0.0069

Saved model and scalers: lstm_tplus7_stacked

48/48 ————— 1s 11ms/step

Train R²: 0.9700, RMSE: 466.63, MAE: 365.38



Temporal Convolutional Networks

```
In [15]: Dropped_tcn_cols = [
    'Open', 'High', 'Low', 'Volume', 'Adj Close',
    'Volatility_Log_10', 'cl-op', 'hi-lo', 'pct_change',
    'total_buying_intent', 'total_prediction_intent'
]
df_tcn_filtered = df_tcn.drop(columns=Dropped_tcn_cols).copy()
```

```

In [24]: def train_tcn_logreturn_model(df_tcn_filtered, df_targets, forecast_horizon
        =1, window_size=30, epochs=50, batch_size=32):
        print(f"\nTraining TCN model to predict Log Returns at t+{forecast_hori
        zon}")

        target_series = df_tcn_filtered['Log_Returns'].shift(-forecast_hori
        zon).dropna()
        df_inputs = df_tcn_filtered.iloc[:len(target_series)]

        # Scaling
        X_scaler = MinMaxScaler()
        y_scaler = MinMaxScaler()
        X_scaled = X_scaler.fit_transform(df_inputs)
        y_scaled = y_scaler.fit_transform(target_series.values.reshape(-1, 1))

        # Sequences
        X_seq, y_seq = [], []
        for i in range(window_size, len(X_scaled)):
            X_seq.append(X_scaled[i - window_size:i])
            y_seq.append(y_scaled[i])
        X_seq, y_seq = np.array(X_seq), np.array(y_seq)

        assert len(X_seq) == len(y_seq), "Sequence and target length mismatch"

        # Build model
        model = Sequential()
        model.add(TCN(
            input_shape=(X_seq.shape[1], X_seq.shape[2]),
            nb_filters=32,
            kernel_size=2,
            dropout_rate=0.1,
        ))
        model.add(GroupNormalization(groups=8))
        model.add(Dropout(0.2))
        model.add(Dense(1)) # Linear output
        model.compile(optimizer=Adam(learning_rate=1e-4), loss=Huber())

        early_stop = EarlyStopping(monitor='val_loss', patience=5, restore_best
        _weights=True)

        # Fit model
        print("Starting model.fit()")
        model.fit(
            X_seq, y_seq,
            epochs=epochs,
            batch_size=batch_size,
            validation_split=0.1,
            callbacks=[early_stop],
            verbose=2
        )
        print("model.fit() complete")

        # Save model & scalers
        model_name = f"tcn_logret_tplus{forecast_horizon}"
        base_path = "B:/DCU/Practicum/Proj/Models"
        os.makedirs(base_path, exist_ok=True)
        model.save(f"{base_path}/{model_name}.keras")
        joblib.dump(X_scaler, f"{base_path}/{model_name}_scalerX.pkl")
        joblib.dump(y_scaler, f"{base_path}/{model_name}_scalerY.pkl")
        joblib.dump(df_tcn_filtered.columns.tolist(), f"{base_path}/{model_nam

```

```

e}_features.pkl")
    print(f"Saved model, scalers, and features for {model_name}")

    # Predict
    y_pred_scaled = model.predict(X_seq)
    y_pred_log = y_scaler.inverse_transform(y_pred_scaled).flatten()

    # clip predicted log returns to avoid exponential blowups
    y_pred_log = np.clip(y_pred_log, -0.15, 0.15)

    # Reconstruct close price
    close_start_idx = window_size - 1
    close_end_idx = close_start_idx + len(y_pred_log)
    close_t = df_targets['Close'].iloc[close_start_idx:close_end_idx].values

    assert len(close_t) == len(y_pred_log), "Mismatch: Close_t vs Predicted log returns"
    y_pred_close = close_t * np.exp(y_pred_log)

    # Actual Close at t+h
    y_true_close = df_targets['Close'].shift(-forecast_horizon).dropna().iloc[window_size:]
    y_true_close = y_true_close[:len(y_pred_close)].values

    date_series = multimodal['Date'].iloc[window_size + forecast_horizon: window_size + forecast_horizon + len(y_pred_close)]

    # Evaluation
    r2 = r2_score(y_true_close, y_pred_close)
    rmse = np.sqrt(mean_squared_error(y_true_close, y_pred_close))
    mae = mean_absolute_error(y_true_close, y_pred_close)
    print(f"TCN Train R²: {r2:.4f}, RMSE: {rmse:.2f}, MAE: {mae:.2f}")

    # Save metrics
    metrics_path = os.path.join(base_path, f"{model_name}_metrics.txt")
    with open(metrics_path, "w") as f:
        f.write(f"Forecast Horizon = t+{forecast_horizon}\n")
        f.write(f"Train R² = {r2:.4f}\n")
        f.write(f"Train RMSE = {rmse:.2f}\n")
        f.write(f"Train MAE = {mae:.2f}\n")
    print(f"Saved training metrics to {metrics_path}")

    # Plot
    plt.figure(figsize=(14, 5))
    plt.plot(date_series, y_true_close, label='Actual', color='blue')
    plt.plot(date_series, y_pred_close, label='Predicted', color='orange')
    plt.title(f"TCN Forecast (t+{forecast_horizon}) - Training Set")
    plt.xlabel("Date")
    plt.ylabel("Close Price")
    plt.legend()
    plt.tight_layout()
    plt.savefig(f"B:/DCU/Practicum/Proj/Outputs/tcn_train_t_plus_{forecast_horizon}.png")
    plt.show()

    return model, {"r2": r2, "rmse": rmse, "mae": mae}

for h in [1, 3, 7]:
    clear_tf_memory()

```

```
train_tcn_logreturn_model(df_tcn_filtered, df_targets, forecast_horizon  
=h)
```

```
Training TCN model to predict Log Returns at t+1
Starting model.fit()
Epoch 1/50
44/44 - 4s - 102ms/step - loss: 0.4008 - val_loss: 0.0530
Epoch 2/50
44/44 - 0s - 11ms/step - loss: 0.3013 - val_loss: 0.0229
Epoch 3/50
44/44 - 0s - 11ms/step - loss: 0.2594 - val_loss: 0.0144
Epoch 4/50
44/44 - 0s - 11ms/step - loss: 0.2413 - val_loss: 0.0122
Epoch 5/50
44/44 - 0s - 11ms/step - loss: 0.2120 - val_loss: 0.0110
Epoch 6/50
44/44 - 0s - 11ms/step - loss: 0.2001 - val_loss: 0.0075
Epoch 7/50
44/44 - 0s - 11ms/step - loss: 0.1776 - val_loss: 0.0055
Epoch 8/50
44/44 - 0s - 11ms/step - loss: 0.1722 - val_loss: 0.0054
Epoch 9/50
44/44 - 0s - 10ms/step - loss: 0.1519 - val_loss: 0.0056
Epoch 10/50
44/44 - 0s - 11ms/step - loss: 0.1408 - val_loss: 0.0034
Epoch 11/50
44/44 - 0s - 11ms/step - loss: 0.1413 - val_loss: 0.0021
Epoch 12/50
44/44 - 0s - 11ms/step - loss: 0.1349 - val_loss: 0.0019
Epoch 13/50
44/44 - 0s - 10ms/step - loss: 0.1249 - val_loss: 0.0018
Epoch 14/50
44/44 - 0s - 10ms/step - loss: 0.1233 - val_loss: 0.0016
Epoch 15/50
44/44 - 0s - 10ms/step - loss: 0.1097 - val_loss: 0.0017
Epoch 16/50
44/44 - 0s - 10ms/step - loss: 0.1098 - val_loss: 0.0014
Epoch 17/50
44/44 - 0s - 10ms/step - loss: 0.1096 - val_loss: 0.0014
Epoch 18/50
44/44 - 0s - 10ms/step - loss: 0.1030 - val_loss: 0.0017
Epoch 19/50
44/44 - 0s - 10ms/step - loss: 0.1002 - val_loss: 0.0015
Epoch 20/50
44/44 - 0s - 10ms/step - loss: 0.0982 - val_loss: 0.0013
Epoch 21/50
44/44 - 0s - 10ms/step - loss: 0.0967 - val_loss: 0.0015
Epoch 22/50
44/44 - 0s - 10ms/step - loss: 0.0982 - val_loss: 0.0011
Epoch 23/50
44/44 - 0s - 10ms/step - loss: 0.0958 - val_loss: 0.0011
Epoch 24/50
44/44 - 0s - 11ms/step - loss: 0.0873 - val_loss: 0.0011
Epoch 25/50
44/44 - 0s - 11ms/step - loss: 0.0869 - val_loss: 0.0011
Epoch 26/50
44/44 - 0s - 10ms/step - loss: 0.0836 - val_loss: 0.0010
Epoch 27/50
44/44 - 0s - 10ms/step - loss: 0.0814 - val_loss: 0.0011
Epoch 28/50
44/44 - 0s - 10ms/step - loss: 0.0813 - val_loss: 0.0011
Epoch 29/50
44/44 - 0s - 10ms/step - loss: 0.0747 - val_loss: 9.8430e-04
Epoch 30/50
```

44/44 - 0s - 10ms/step - loss: 0.0721 - val_loss: 9.4004e-04

Epoch 31/50

44/44 - 0s - 10ms/step - loss: 0.0729 - val_loss: 0.0011

Epoch 32/50

44/44 - 0s - 10ms/step - loss: 0.0704 - val_loss: 0.0011

Epoch 33/50

44/44 - 0s - 10ms/step - loss: 0.0657 - val_loss: 0.0013

Epoch 34/50

44/44 - 0s - 10ms/step - loss: 0.0632 - val_loss: 9.6007e-04

Epoch 35/50

44/44 - 0s - 10ms/step - loss: 0.0591 - val_loss: 0.0010


model.fit() complete

Saved model, scalers, and features for tcn_logret_tplus1

49/49 ————— **1s** 9ms/step

TCN Train R²: 0.9950, RMSE: 192.23, MAE: 141.85

Saved training metrics to B:/DCU/Practicum/Proj/Models\tcn_logret_tplus1_metrics.txt



Training TCN model to predict Log Returns at $t+3$

Starting model.fit()

Epoch 1/50

44/44 - 4s - 98ms/step - loss: 0.5472 - val_loss: 0.0603

Epoch 2/50

44/44 - 0s - 10ms/step - loss: 0.4133 - val_loss: 0.0359

Epoch 3/50

44/44 - 0s - 10ms/step - loss: 0.3419 - val_loss: 0.0247

Epoch 4/50

44/44 - 0s - 10ms/step - loss: 0.2831 - val_loss: 0.0216

Epoch 5/50

44/44 - 0s - 10ms/step - loss: 0.2726 - val_loss: 0.0148

Epoch 6/50

44/44 - 0s - 10ms/step - loss: 0.2460 - val_loss: 0.0125

Epoch 7/50

44/44 - 0s - 10ms/step - loss: 0.2292 - val_loss: 0.0067

Epoch 8/50

44/44 - 0s - 10ms/step - loss: 0.2091 - val_loss: 0.0078

Epoch 9/50

44/44 - 0s - 10ms/step - loss: 0.2032 - val_loss: 0.0077

Epoch 10/50

44/44 - 0s - 10ms/step - loss: 0.1896 - val_loss: 0.0093

Epoch 11/50

44/44 - 0s - 10ms/step - loss: 0.1929 - val_loss: 0.0102

Epoch 12/50

44/44 - 0s - 10ms/step - loss: 0.1851 - val_loss: 0.0091


model.fit() complete

Saved model, scalers, and features for tcn_logret_tplus3

49/49 ————— 1s 9ms/step

TCN Train R^2 : 0.9845, RMSE: 337.50, MAE: 261.15

Saved training metrics to B:/DCU/Practicum/Proj/Models\tcn_logret_tplus3_metrics.txt




```
Training TCN model to predict Log Returns at t+7
Starting model.fit()
Epoch 1/50
44/44 - 4s - 99ms/step - loss: 0.4403 - val_loss: 0.1279
Epoch 2/50
44/44 - 0s - 10ms/step - loss: 0.3425 - val_loss: 0.0485
Epoch 3/50
44/44 - 0s - 10ms/step - loss: 0.3295 - val_loss: 0.0324
Epoch 4/50
44/44 - 0s - 11ms/step - loss: 0.2898 - val_loss: 0.0197
Epoch 5/50
44/44 - 0s - 10ms/step - loss: 0.2644 - val_loss: 0.0158
Epoch 6/50
44/44 - 0s - 10ms/step - loss: 0.2356 - val_loss: 0.0181
Epoch 7/50
44/44 - 0s - 10ms/step - loss: 0.2451 - val_loss: 0.0156
Epoch 8/50
44/44 - 0s - 10ms/step - loss: 0.2230 - val_loss: 0.0088
Epoch 9/50
44/44 - 0s - 10ms/step - loss: 0.2096 - val_loss: 0.0107
Epoch 10/50
44/44 - 0s - 11ms/step - loss: 0.1912 - val_loss: 0.0031
Epoch 11/50
44/44 - 0s - 10ms/step - loss: 0.1820 - val_loss: 0.0033
Epoch 12/50
44/44 - 0s - 10ms/step - loss: 0.1708 - val_loss: 0.0033
Epoch 13/50
44/44 - 0s - 10ms/step - loss: 0.1656 - val_loss: 0.0019
Epoch 14/50
44/44 - 0s - 11ms/step - loss: 0.1651 - val_loss: 0.0019
Epoch 15/50
44/44 - 0s - 10ms/step - loss: 0.1495 - val_loss: 0.0016
Epoch 16/50
44/44 - 0s - 10ms/step - loss: 0.1520 - val_loss: 0.0014
Epoch 17/50
44/44 - 1s - 12ms/step - loss: 0.1432 - val_loss: 0.0020
Epoch 18/50
44/44 - 1s - 13ms/step - loss: 0.1260 - val_loss: 0.0014
Epoch 19/50
44/44 - 0s - 11ms/step - loss: 0.1200 - val_loss: 0.0012
Epoch 20/50
44/44 - 0s - 10ms/step - loss: 0.1265 - val_loss: 0.0012
Epoch 21/50
44/44 - 1s - 12ms/step - loss: 0.1196 - val_loss: 0.0026
Epoch 22/50
44/44 - 1s - 12ms/step - loss: 0.1076 - val_loss: 9.2440e-04
Epoch 23/50
44/44 - 0s - 10ms/step - loss: 0.1013 - val_loss: 9.5243e-04
Epoch 24/50
44/44 - 0s - 11ms/step - loss: 0.1033 - val_loss: 8.7182e-04
Epoch 25/50
44/44 - 0s - 11ms/step - loss: 0.0967 - val_loss: 0.0014
Epoch 26/50
44/44 - 0s - 11ms/step - loss: 0.0955 - val_loss: 0.0015
Epoch 27/50
44/44 - 1s - 13ms/step - loss: 0.0878 - val_loss: 0.0015
Epoch 28/50
44/44 - 1s - 12ms/step - loss: 0.0882 - val_loss: 8.0366e-04
Epoch 29/50
44/44 - 1s - 11ms/step - loss: 0.0872 - val_loss: 9.2579e-04
Epoch 30/50
```



```


44/44 - 1s - 13ms/step - loss: 0.0888 - val_loss: 7.8006e-04
Epoch 31/50
44/44 - 1s - 13ms/step - loss: 0.0786 - val_loss: 8.0489e-04
Epoch 32/50
44/44 - 0s - 11ms/step - loss: 0.0810 - val_loss: 7.1540e-04
Epoch 33/50
44/44 - 0s - 11ms/step - loss: 0.0760 - val_loss: 6.7770e-04
Epoch 34/50
44/44 - 0s - 11ms/step - loss: 0.0715 - val_loss: 6.5881e-04
Epoch 35/50
44/44 - 0s - 10ms/step - loss: 0.0740 - val_loss: 6.8288e-04
Epoch 36/50
44/44 - 1s - 14ms/step - loss: 0.0650 - val_loss: 6.4124e-04
Epoch 37/50
44/44 - 1s - 14ms/step - loss: 0.0655 - val_loss: 6.3827e-04
Epoch 38/50
44/44 - 1s - 13ms/step - loss: 0.0652 - val_loss: 6.7055e-04
Epoch 39/50
44/44 - 1s - 13ms/step - loss: 0.0609 - val_loss: 6.8099e-04
Epoch 40/50
44/44 - 1s - 14ms/step - loss: 0.0651 - val_loss: 7.0217e-04
Epoch 41/50
44/44 - 1s - 13ms/step - loss: 0.0552 - val_loss: 6.5736e-04
Epoch 42/50
44/44 - 1s - 12ms/step - loss: 0.0542 - val_loss: 7.7907e-04
model.fit() complete
Saved model, scalers, and features for tcn_logret_tplus7
49/49 ————— 1s 9ms/step
TCN Train R2: 0.9835, RMSE: 348.33, MAE: 260.38
Saved training metrics to B:/DCU/Practicum/Proj/Models\tcn_logret_tplus7_m
etrics.txt

```



```
In [17]: sns.histplot(df_tcn_filtered['Log_Returns'], bins=100, kde=True)
```

```
Out[17]: <Axes: xlabel='Log_Returns', ylabel='Count'>
```



```
In [18]: print("Skew:", skew(df_tcn_filtered['Log_Returns']))  
print("Kurtosis:", kurtosis(df_tcn_filtered['Log_Returns']))
```

Skew: 0.06939321499881877

Kurtosis: 10.324332846429078

```

In [19]: def run_shap_kernel_on_tcn(model_dir, df_input, forecast_horizon, window_size=30, num_samples=25):
    model_name = f"tcn_logret_tplus{forecast_horizon}"
    model_path = os.path.join(model_dir, f"{model_name}.keras")
    scalerX_path = os.path.join(model_dir, f"{model_name}_scalerX.pkl")

    # Load model and scaler
    model = load_model(model_path, compile=False)
    X_scaler = joblib.load(scalerX_path)

    # Scale and sequence input
    df_scaled = X_scaler.transform(df_input)
    X_seq = []
    for i in range(window_size, len(df_scaled)):
        X_seq.append(df_scaled[i - window_size:i])
    X_seq = np.array(X_seq)

    # Sample last N sequences
    X_sampled = X_seq[-num_samples:]

    # Track shape dimensions
    num_samples_actual = X_sampled.shape[0]
    num_timesteps = X_sampled.shape[1]
    num_features = X_sampled.shape[2]

    # Flatten input for KernelExplainer
    X_flat = X_sampled.reshape((num_samples_actual, num_timesteps * num_features))

    # Define prediction wrapper
    def predict_fn(x_flat):
        x_resaped = x_flat.reshape((-1, num_timesteps, num_features))
        preds = model.predict(x_resaped)
        return np.array(preds).astype(np.float64).reshape(-1, 1)

    # Run SHAP KernelExplainer
    explainer = shap.KernelExplainer(predict_fn, X_flat)
    try:
        shap_values = explainer(X_flat)
        with open(f"{model_name}_shap_dump.pkl", "wb") as f:
            joblib.dump(shap_values, f)
    except Exception as e:
        print(f"SHAP failed: {e}")
        return

    shap_array = np.abs(shap_values.values) # shape: [samples, window_size * features]
    shap_array_3d = shap_array.reshape(shap_array.shape[0], window_size, num_features)
    shap_feature_mean = np.mean(shap_array_3d, axis=(0, 1)) # average across time and samples

    feature_names = df_input.columns.tolist()

    # Rank + plot
    shap_df = list(zip(feature_names, shap_feature_mean))
    shap_df.sort(key=lambda x: x[1], reverse=True)
    sorted_features, sorted_importance = zip(*shap_df)

    plt.figure(figsize=(10, 6))



```




```
plt.barh(sorted_features[::-1], sorted_importance[::-1], color='steelblue')
plt.title(f"SHAP Feature Importance – TCN KernelExplainer (t+{forecast_
horizon})")
plt.tight_layout()
plt.xlabel("Mean SHAP Value")
plt.ylabel("Features")
plt.savefig(f"B:/DCU/Practicum/Proj/Outputs/{model_name}_shap_kernel.png")
plt.show()




model_dir = "B:/DCU/Practicum/Proj/Models"




# for h in [1, 3, 7]:
#     print(f"\n Running SHAP KernelExplainer for TCN t+{h}")
#     run_shap_kernel_on_tcn(model_dir, df_tcn, forecast_horizon=h, num_
samples=100)




run_shap_kernel_on_tcn(model_dir, df_tcn_filtered, forecast_horizon=1, num_
samples=25)
```




1/1  0s 271ms/step
0%|  | 0/25 [00:00<?, ?it/s]




1/1  0s 288ms/step
3429/3429  11s 3ms/step
4%|  | 1/25 [00:13<05:18, 13.26s/it]




1/1  0s 41ms/step
3429/3429  11s 3ms/step
8%|  | 2/25 [00:25<04:57, 12.94s/it]




1/1  0s 30ms/step
3429/3429  11s 3ms/step
12%|  | 3/25 [00:38<04:42, 12.84s/it]




1/1  0s 26ms/step
3429/3429  11s 3ms/step
16%|  | 4/25 [00:51<04:31, 12.95s/it]




1/1  0s 27ms/step
3429/3429  11s 3ms/step
20%|  | 5/25 [01:04<04:17, 12.85s/it]




1/1  0s 27ms/step
3429/3429  11s 3ms/step
24%|  | 6/25 [01:17<04:02, 12.79s/it]




1/1  0s 29ms/step
3429/3429  11s 3ms/step
28%|  | 7/25 [01:29<03:48, 12.70s/it]



1/1  0s 27ms/step
3429/3429  11s 3ms/step
32%|  | 8/25 [01:42<03:35, 12.66s/it]

1/1  0s 26ms/step
3429/3429  11s 3ms/step
36%|  | 9/25 [01:54<03:21, 12.57s/it]

1/1  0s 28ms/step
3429/3429  11s 3ms/step
40%|  | 10/25 [02:07<03:08, 12.59s/it]

1/1  0s 34ms/step
3429/3429  13s 4ms/step
44%|  | 11/25 [02:22<03:08, 13.44s/it]

1/1  0s 30ms/step
3429/3429  11s 3ms/step
48%|  | 12/25 [02:35<02:53, 13.35s/it]

1/1  0s 28ms/step
3429/3429  11s 3ms/step

52%|██████| | 13/25 [02:48<02:38, 13.22s/it]
1/1 ██████████ 0s 30ms/step
3429/3429 ██████████ 11s 3ms/step

56%|██████| | 14/25 [03:01<02:24, 13.15s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

60%|██████| | 15/25 [03:14<02:10, 13.06s/it]
1/1 ██████████ 0s 27ms/step
3429/3429 ██████████ 11s 3ms/step

64%|██████| | 16/25 [03:27<01:56, 12.98s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

68%|██████| | 17/25 [03:40<01:43, 12.93s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

72%|██████| | 18/25 [03:52<01:30, 12.87s/it]
1/1 ██████████ 0s 27ms/step
3429/3429 ██████████ 11s 3ms/step

76%|██████| | 19/25 [04:05<01:17, 12.89s/it]
1/1 ██████████ 0s 29ms/step
3429/3429 ██████████ 11s 3ms/step

80%|██████| | 20/25 [04:18<01:04, 12.88s/it]
1/1 ██████████ 0s 27ms/step
3429/3429 ██████████ 11s 3ms/step


84%|██████| | 21/25 [04:31<00:51, 12.84s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

88%|██████| | 22/25 [04:44<00:38, 12.79s/it]
1/1 ██████████ 0s 31ms/step
3429/3429 ██████████ 11s 3ms/step



92%|██████| | 23/25 [04:56<00:25, 12.80s/it]
1/1 ██████████ 0s 27ms/step
3429/3429 ██████████ 11s 3ms/step




96%|██████| | 24/25 [05:10<00:13, 13.02s/it]
1/1 ██████████ 0s 29ms/step
3429/3429 ██████████ 11s 3ms/step




100%|██████| | 25/25 [05:23<00:00, 12.94s/it]









```
In [25]: run_shap_kernel_on_tcn(model_dir, df_tcn_filtered, forecast_horizon=3, num_
samples=25)
```





1/1  0s 293ms/step
0%|  | 0/25 [00:00<?, ?it/s]




1/1  0s 280ms/step
3429/3429  13s 4ms/step
4%|  | 1/25 [00:15<06:05, 15.21s/it]




1/1  0s 37ms/step
3429/3429  12s 3ms/step
8%|  | 2/25 [00:29<05:32, 14.47s/it]




1/1  0s 32ms/step
3429/3429  12s 3ms/step
12%|  | 3/25 [00:42<05:09, 14.05s/it]




1/1  0s 28ms/step
3429/3429  12s 3ms/step
16%|  | 4/25 [00:56<04:51, 13.90s/it]




1/1  0s 28ms/step
3429/3429  12s 3ms/step
20%|  | 5/25 [01:09<04:34, 13.74s/it]




1/1  0s 29ms/step
3429/3429  12s 3ms/step
24%|  | 6/25 [01:23<04:19, 13.65s/it]




1/1  0s 28ms/step
3429/3429  12s 3ms/step
28%|  | 7/25 [01:36<04:04, 13.57s/it]



1/1  0s 27ms/step
3429/3429  12s 3ms/step
32%|  | 8/25 [01:50<03:49, 13.50s/it]

1/1  0s 28ms/step
3429/3429  12s 3ms/step
36%|  | 9/25 [02:03<03:35, 13.45s/it]

1/1  0s 33ms/step
3429/3429  11s 3ms/step
40%|  | 10/25 [02:16<03:20, 13.36s/it]

1/1  0s 27ms/step
3429/3429  11s 3ms/step
44%|  | 11/25 [02:29<03:06, 13.31s/it]

1/1  0s 28ms/step
3429/3429  11s 3ms/step
48%|  | 12/25 [02:43<02:53, 13.31s/it]

1/1  0s 29ms/step
3429/3429  11s 3ms/step

52%|██████| | 13/25 [02:56<02:39, 13.30s/it]
1/1 ————— 0s 29ms/step
3429/3429 ————— 11s 3ms/step

56%|██████| | 14/25 [03:09<02:25, 13.26s/it]
1/1 ————— 0s 35ms/step
3429/3429 ————— 12s 3ms/step

60%|██████| | 15/25 [03:22<02:12, 13.29s/it]
1/1 ————— 0s 27ms/step
3429/3429 ————— 12s 3ms/step

64%|██████| | 16/25 [03:36<02:00, 13.37s/it]
1/1 ————— 0s 28ms/step
3429/3429 ————— 11s 3ms/step

68%|██████| | 17/25 [03:49<01:46, 13.33s/it]
1/1 ————— 0s 33ms/step
3429/3429 ————— 12s 3ms/step

72%|██████| | 18/25 [04:03<01:33, 13.39s/it]
1/1 ————— 0s 34ms/step
3429/3429 ————— 11s 3ms/step

76%|██████| | 19/25 [04:16<01:20, 13.38s/it]
1/1 ————— 0s 28ms/step
3429/3429 ————— 12s 3ms/step

80%|██████| | 20/25 [04:30<01:07, 13.40s/it]
1/1 ————— 0s 31ms/step
3429/3429 ————— 11s 3ms/step


84%|██████| | 21/25 [04:43<00:53, 13.37s/it]
1/1 ————— 0s 29ms/step
3429/3429 ————— 12s 3ms/step

88%|██████| | 22/25 [04:56<00:40, 13.44s/it]
1/1 ————— 0s 48ms/step
3429/3429 ————— 13s 4ms/step



92%|██████| | 23/25 [05:11<00:27, 13.78s/it]
1/1 ————— 0s 43ms/step
3429/3429 ————— 13s 4ms/step




96%|██████| | 24/25 [05:26<00:14, 14.04s/it]
1/1 ————— 0s 40ms/step
3429/3429 ————— 13s 4ms/step




100%|██████| | 25/25 [05:41<00:00, 13.64s/it]









```
In [26]: run_shap_kernel_on_tcn(model_dir, df_tcn_filtered, forecast_horizon=7, num_
         samples=25)
```




1/1  0s 277ms/step
0%|  | 0/25 [00:00<?, ?it/s]




1/1  0s 294ms/step
3429/3429  13s 4ms/step
4%|  | 1/25 [00:14<05:54, 14.75s/it]




1/1  0s 38ms/step
3429/3429  13s 4ms/step
8%|  | 2/25 [00:29<05:37, 14.70s/it]




1/1  0s 31ms/step
3429/3429  13s 4ms/step
12%|  | 3/25 [00:43<05:20, 14.58s/it]




1/1  0s 30ms/step
3429/3429  13s 4ms/step
16%|  | 4/25 [00:58<05:05, 14.52s/it]




1/1  0s 30ms/step
3429/3429  12s 4ms/step
20%|  | 5/25 [01:12<04:49, 14.49s/it]




1/1  0s 35ms/step
3429/3429  12s 4ms/step
24%|  | 6/25 [01:27<04:34, 14.45s/it]




1/1  0s 60ms/step
3429/3429  13s 4ms/step
28%|  | 7/25 [01:41<04:20, 14.45s/it]



1/1  0s 31ms/step
3429/3429  13s 4ms/step
32%|  | 8/25 [01:55<04:05, 14.42s/it]

1/1  0s 30ms/step
3429/3429  12s 4ms/step
36%|  | 9/25 [02:10<03:50, 14.38s/it]

1/1  0s 38ms/step
3429/3429  13s 4ms/step
40%|  | 10/25 [02:24<03:35, 14.38s/it]

1/1  0s 29ms/step
3429/3429  13s 4ms/step
44%|  | 11/25 [02:38<03:21, 14.36s/it]

1/1  0s 29ms/step
3429/3429  13s 4ms/step
48%|  | 12/25 [02:53<03:07, 14.45s/it]

1/1  0s 31ms/step
3429/3429  13s 4ms/step

52%|███████| | 13/25 [03:07<02:53, 14.44s/it]
1/1 ██████████ 0s 31ms/step
3429/3429 ██████████ 13s 4ms/step

56%|███████| | 14/25 [03:22<02:39, 14.49s/it]
1/1 ██████████ 0s 31ms/step
3429/3429 ██████████ 13s 4ms/step

60%|███████| | 15/25 [03:37<02:25, 14.51s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 13s 4ms/step

64%|███████| | 16/25 [03:52<02:12, 14.74s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

68%|███████| | 17/25 [04:05<01:53, 14.19s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

72%|███████| | 18/25 [04:18<01:36, 13.82s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

76%|███████| | 19/25 [04:31<01:21, 13.53s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

80%|███████| | 20/25 [04:44<01:06, 13.37s/it]
1/1 ██████████ 0s 32ms/step
3429/3429 ██████████ 11s 3ms/step

84%|███████| | 21/25 [04:57<00:52, 13.24s/it]
1/1 ██████████ 0s 33ms/step
3429/3429 ██████████ 11s 3ms/step

88%|███████| | 22/25 [05:09<00:39, 13.14s/it]
1/1 ██████████ 0s 29ms/step
3429/3429 ██████████ 11s 3ms/step

92%|███████| | 23/25 [05:22<00:26, 13.08s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

96%|███████| | 24/25 [05:35<00:13, 13.02s/it]
1/1 ██████████ 0s 28ms/step
3429/3429 ██████████ 11s 3ms/step

100%|███████| | 25/25 [05:48<00:00, 13.95s/it]

